

Exploring a Synthetic SSc Cohort

Data Understanding, Application, and Data Quality

Semion Leyn

2026-09-09

Agenda


- Understanding the dataset
- The application & visualizations
- Data-quality issues and analytical findings
- Design decisions and tradeoffs
- Future work & discussion

Understanding the Dataset

What's in the box

- 11 CSVs, ~4 MB, synthetic and non-PHI
- A patient/subject identifier under **4 different column names** across files — normalized to one canonical `subject_id` at ingestion
- The **Registry**: `demographics` + `ssc_subtype`, a 1:1, 1,500-row master list — every row `study=SSC-REG`, `diagnosis=SSc`
- Every other table: a **partial-coverage** table of clinical/molecular events for a subset of those 1,500 patients

Population and table coverage



A hidden second cohort

4 subject IDs (`SSC_NORM_0101/0102/0104/0110`) appear in `libraries.csv` — but **never** in the Registry.

- `libraries.csv`'s own comments confirm these are “healthy control sample[s]”
- Modeled explicitly as an attribute (`ssc_patient` vs `control`)
- Not a data error
- The QC key-linkage check automates this: confirms all 4 against the known set, and would flag any *future* out-of-Registry `subject_id`.

Static vs. longitudinal, and a naming trap

- `lab_report`, `vitals`, `mrss` — every row reduces to one (`subject`, `date`, `value`) triple (which measure varies per row for labs/vitals; MRSS has only one possible measure). Same shape, unified as one **Observation** abstraction in code (Lab Result / Vital Sign / MRSS stay distinct terms in any user-facing view)
- `bal`, `medications`, `skin_biopsies` — structured event records with no single value: `bal` has two volumes + site + comment, `skin_biopsies` has no numeric value at all (site/indication/specimen/ image path), `medications`

The Application

Streamlit, 4 tabs, SQLite-backed

- **Data & Dictionary** — schema, coverage, QC report link
- **Cohort Overview** — population-level distributions
- **Compare & Discover** — free-form variable picker across any table/level, with a Control Subject overlay toggle
- **Patient Trajectory** — one patient's longitudinal record

Login + at-rest encryption + audit logging gated behind a published demo account (details in Design Decisions).

Data & Dictionary

Data & Dictionary

Cohort Overview

Compare & Discover

Patient Trajectory

QC Report

Logged in as demo

Data & Dictionary

Every table in the built Subject/Cohort store, with live row/column counts, a plain-language description of each field informed by the QC investigation, and up to 3 real example values pulled live from the store.

▼ subjects — 1,504 rows x 2 columns

Every Subject referenced anywhere in the dataset — the supertype spanning both SSc Patients and the 4 Control Subjects.

field	description	examples
subject_id	Canonical Subject identifier, normalized at ingestion.	SSC_NORM_0101, SSC_NORM_0102, SSC_NORM_0104
cohort	'ssc_patient' or 'control' — which Cohort this Subject belongs to.	control, ssc_patient

> demographics — 1,500 rows x 12 columns

> ssc_subtype — 1,500 rows x 7 columns

> vitals — 40,250 rows x 4 columns

> lab_report — 30,941 rows x 4 columns

> mras — 3,113 rows x 4 columns

> pft — 5,069 rows x 5 columns

> antibodies — 3,511 rows x 4 columns

> medications — 3,273 rows x 4 columns

> bal — 1,170 rows x 7 columns

> libraries — 918 rows x 24 columns

> skin_biopsies — 215 rows x 8 columns

Data & Dictionary tab: every table's row/column counts, plain-language field descriptions, and live example values

Cohort Overview

Data & Dictionary

Cohort Overview

Compare & Discover

Patient Trajectory

QC Report

Logged in as demo

Cohort Overview

Population-level structure across the Registry and its 4 Control Subjects — demographic distributions, SSc subtype breakdown, and how many Subjects have at least one record in each clinical table.

Registry patients (ssc_patient)


1500

Control Subjects

4

SSc subtype

dcSSc / lcSSc split (n=1,500 Registry patients)



Cohort Overview tab

Compare & Discover

Data & Dictionary

Cohort Overview

Compare & Discover

Patient Trajectory

QC Report

Logged in as demo

Compare & Discover

Add one comparison panel per pair of variables you want to see — pick an X and a Y independently in each panel, rotate to swap them, and get the chart type that fits what you picked automatically: a scatter for two numeric measures, a box plot for a numeric measure grouped by a category, or a heatmap for two categorical fields.

Panel 1

Remove

X variable

Lab Result: RDW (per-patient summary)

Y variable

Disease Classification: other_dx

Rotate

Color by (optional)

-- select a variable --

subject_id	Lab Result: RDW (per-patient summary)	Disease Classification: other_dx	Cohort
subject_2005	15.9667	ILD; GERD	ssc_patient
subject_2028	12.55	GERD	ssc_patient
subject_2061	15.55	ILD	ssc_patient
subject_2064	15.4	ILD	ssc_patient
subject_2092	12.65	GERD; PAH	ssc_patient
subject_2095	16.2	ILD; GERD	ssc_patient
subject_2099	16	ILD	ssc_patient
subject_2135	12.575	PAH	ssc_patient
subject_2148	15.6	GERD	ssc_patient
subject_2170	15.1	ILD	ssc_patient

Add comparison panel

Compare & Discover tab

Patient Trajectory

Data & Dictionary

Cohort Overview

Compare & Discover

Patient Trajectory

QC Report

Logged in as demo

Subject Filters

Selected values within one filter combine with OR (e.g. picking 2 genders shows Subjects matching either); separate filters combine with AND (e.g. a gender filter plus BAL performed narrows to Subjects matching both).

> Demographics (7 filters)

Time

Visit date

2014-10-272018-07-13

Medication

Ever prescribed

Choose options

Antibody Test

Test result

Choose options

BAL

☐ BAL performed

> Lab Report (19 filters)

Patient Trajectory

Select up to 8 Subjects by `subject_id` to see their longitudinal records overlaid — labs, vitals, MRSS, PFT, and medications plotted over time. Patient name and birth date are never shown here or anywhere else in this app.

subject_id

subject_2064

subject_2092

subject_2135

subject_2151

subject_2160

X-axis

☐ Calendar date


☒ Disease duration (years since onset — first non-Raynaud symptom)

Subject overview


Subject	SSc subtype	Comorbidities	Antibodies	Disease onset	Disease duration (years)
subject_2064	dcSSc	ILD	anti-centromere antibodies: negative; rna polymerase III: negative; scl70: negative	2015-05-04	11.4
subject_2092	lcSSc	GERD; PAH	anti-centromere antibodies: positive; rna polymerase III: negative	2014-12-09	11.8
subject_2135	lcSSc	PAH	anti-centromere antibodies: positive; rna polymerase III: negative	2014-08-20	12.1
subject_2151	dcSSc	--	anti-centromere antibodies: negative; scl70: positive	2015-12-05	10.8
subject_2160	dcSSc	--	scl70: positive	2013-12-18	12.7

Labs


ABSOLUTE BASOPHILS




ABSOLUTE EOSINOPHILS



ABSOLUTE LYMPHOCYTES




ABSOLUTE MONOCYTE



Patient Trajectory tab

Data-Quality Issues and Analytical Findings

Height recorded in mixed units




1,235 of 1,500 patients in **inches** (54.7-78.5); the other 265 in **centimeters** (141.3-194.7 cm). Fixed at ingestion: cm-scale values converted to inches, raw CSV left untouched.

One high-confidence data-entry error

`demographics.csv` lists `subject_8545`'s weight as **22.5** (impossible for a living adult) — but 5 separate `vitals.csv` visits spanning 2017-2020 all report ~160.6 lbs for the same patient.

Decision: flagged in the QC report, not silently corrected. The vitals-based value is well-evidenced, but `demographics.weight` is a distinct static field with no documented derivation from vitals — overwriting it would assert a number we don't have provenance for.

A synthetic-data placeholder, not a clinical signal




WEIGHT IN POUND = 160.6 appears 870 times across 195 distinct patients — the next most common value appears only 39 times. Some patients show this exact

The same pattern in BP and pulse, measured differently

The per-patient check above is blind to a filler used on only *some* of a patient's visits. Looking at each value's share of *every* record instead:

Measure	Value	Share of all records	vs. local background
BP SYSTOLIC	124	25% of 7,848	11.7x
BP DIASTOLIC	77	25% of 7,848	5.7x
PULSE	76	16% of 8,381	4.6x

PFT bounds look like generator clipping



Three physiologically unrelated measures all range **exactly 40.0-130.0% predicted**. A clean guideline-range check here is itself the finding — reported as a

The same generic check finds real injected errors

Running `check_guideline_range` against `lab_report.csv` for the first time surfaces isolated, physiologically-impossible values nobody had manually audited:

- `ABSOLUTE LYMPHOCYTES`, `MCH`, `MCHC` — one `9999` sentinel each
- `WBC` — two `999` sentinels
- `HEMOGLOBIN`, `NEUTROPHILS`, `LYMPHOCYTES`, `MCV`, `NM BKR`
`ABSOLUTE NEUTROPHIL` — 1-2 negative values each (impossible for any of those measures)

Medications: names fragmented, doses mostly clean

Only 17 distinct medication strings total — but `CellCept` (192), `mycophenolate mofetil` (211), and `MMF` (189) are the same drug under three names. Merged, the total (592) is nearly 2.5x the next most-prescribed drug (rituximab, 253).

- **Dose:** only 24 distinct strings; a few look like intentionally injected errors — negative doses, `999 tablets daily`, `250 g twice daily` (likely a mg/g slip, ~1000x too high), `500 mL tablets daily` (a volume unit

PII-shaped fields in a synthetic dataset

`demographics` has first/last name and birth date;
`mrss/skin_biopsies` have `ENTRY_USER_NAME` (a clinician name) — despite the dataset being confirmed synthetic/non-PHI.

Decision: patient name/DOB suppressed everywhere downstream (`subject_id` + age-at-event only) — a deliberate demonstration of PHI-handling judgment.

`ENTRY_USER_NAME` kept visible: provider-side operational metadata, not what HIPAA governs, and useful for a “does score entry vary by rater?” angle

Design Decisions and Tradeoffs

HIPAA controls, implemented literally

Real login, real encryption at rest, real audit logging — despite the data being confirmed synthetic/non-PHI and the deliverable requiring a public repo + public app link.

Gated by a published demo account, so the controls are **genuine engineering rather than theater**, without blocking reviewer access.

(ADR 0001)

Encrypt the store, not the source CSVs

- Keeps the repo **clone-and-run reproducible** — a reviewer can inspect the actual source data directly
- Still gives a real “encryption at rest” component to explain: the runtime SQLite store the app/notebook actually reads from is encrypted

(ADR 0002)

SQLite with a Subject supertype, not pandas end-to-end

Chosen specifically because it **models the Control Subject / Registry relationship correctly** — cohort membership, not a foreign-key violation to explain away.

The role's own scope ("data modeling across clinical data sources") rewards showing schema design even at a data size where SQL isn't strictly necessary for performance.

(ADR 0005)

Demo secrets committed, against normal practice

The credentials and encryption key are **deliberately public by design** (same reasoning as the HIPAA-literal controls) — keeping them out of git would only add setup friction for reviewers, with no actual security benefit.

Documented explicitly as a demo-only deviation.

(ADR 0004)

Scope triage under a same-day deadline

- Flat package layout over a layered architecture
- Minimal-but-real CI: ruff + mypy + pytest via GitHub Actions
- Targeted tests over exhaustive coverage
- The originally-planned 5th “discovery” app page merged into the group-comparison page once it became clear they were the same free-form explorer

Future Work

Known, deliberately scoped — not bugs. The app is stable, tested, deployed, green.

Addressed since the initial clinical read:

- Disease-duration x-axis (years-since-onset toggle, alongside calendar date) — Patient Trajectory & Cohort Overview
- Patient Trajectory subject header — subtype, comorbidities, antibodies, disease onset/duration
- Compare & Discover “color by” facet — a first step toward subtype/antibody-faceted outcome views

Thank you — discussion

Repo: (public GitHub link) · App: nw-ssc-de.streamlit.app

Data Understanding, Application, and Data Quality

Appendix

Appendix: HIPAA-style controls, in detail

- **Login:** gated by a published demo account ([demo / SscDemo2026!](#)) — a demonstration of the control, not a real access boundary
- **Encryption at rest:** the runtime SQLite store is built encrypted from the source CSVs at startup; source CSVs themselves are left plaintext for reviewer reproducibility
- **Audit logging:** every login attempt (success/failure) appended to a local [audit.log](#) with timestamp and user; does not persist across restarts under the chosen

Appendix: QC methodology detail

`check_constant_value_across_visits` — per-patient dominance ratio: share of patients for whom a value is invariant across *all* their visits, vs. expectation if personal-constant values were spread evenly across observed distinct values. Requires ≥ 2 visits, ≥ 10 subjects.

`check_value_spike` — population-level share of *every* record for a measure landing on one value, vs. the local background set by its nearest numeric neighbors.

Complementary to the check above: catches a filler used on only some of a patient's visits, which the per-patient check is blind to.

`check_guideline_range` — generic, reusable against any

Appendix: Future-work gaps, expanded

Originally 7 gaps flagged from a clinical read of the app ([.scratch/HANDOFF.md](#), “What’s NOT done — real gaps, not polish”). Since then: 3 addressed (disease-duration axis, Patient Trajectory subject header, Compare & Discover color-by), 1 folded into [CONTEXT.md](#) as a documented data limitation, 1 dropped after consideration (PFT decline flagging). The 2 still open (event overlay, medication classing) read as “*what would make this SSc-domain-aware rather than generic-longitudinal-EHR-browser-aware*” to a pulmonary-research committee specifically — not code-quality gaps.